RIGHTSIZE AI

The AI Value Intelligence Platform for Regulated Enterprises

Guillaume Ehinger · Etay Zwick · Alex Paterson · David Mindham · Ravikumaran Govender

TRIUM Global EMBA · NYU Stern / LSE / HEC Paris · Capstone Prospectus · May 2026

THE PROBLEM

What is the problem?

AI spend is compounding. *Nobody can account for it.*

~30%

of AI spend is wasted (BCG)

95%

of AI pilots: no measurable P&L impact (MIT)

Every AI request is a decision. Today, nobody is effectively managing these decisions. As enterprise AI adoption accelerates — average monthly AI spend has reached $85,000 and is growing 36% year-on-year — a structural accountability gap has opened up that no existing tool addresses. For regulated, mid-market enterprises spending c.$500k annually on AI, this implies c.$150k in unmanaged waste, before accounting for regulatory and risk exposure, duplicative tooling or remediation costs from governance failures.

CFO — Financial blind spot

Only 29% of executives can measure AI ROI with confidence. 95% of generative AI pilots deliver no measurable P&L impact (MIT NANDA, 2025). When the board asks 'what are we getting for this?', no one can answer. Teams default to the most powerful - and most expensive - models for every task regardless of need. An estimated 30% of enterprise AI spend is outright waste.

CIO/CTO — Operational chaos

AI adoption has spread well beyond IT. Marketing, finance, product, legal, and HR are all generating AI costs independently, often using unsanctioned tools. Over 90% of organisations have employees using personal AI tools for work (shadow AI). There is no consolidated view of the AI portfolio — no inventory, no tagging, no allocation by use case or business outcome.

CRO — Invisible risk

Proprietary data is reaching unvetted models with no audit trail. Governance frameworks have not kept pace with adoption. The EU AI Act's high-risk obligations take effect August 2026, creating legal exposure with fines up to $15M or 3% of global turnover. The Corporate Sustainability Reporting Directive requires AI energy consumption disclosure. Sector mandates in financial services, insurance, and healthcare are tightening further.

Enterprises intend to scale AI innovation — and instead end up scaling blank-cheque liability.

MARKET OPPORTUNITY

What is the market?

A potential $4.5B accountability gap - with regulatory deadlines accelerating demand

The baseline. Approximately 50,000 mid-market firms globally ($50M-$1B revenue, non-technology sectors) are actively deploying AI. These firms spend an average of $300,000 per year on AI services, representing roughly $15 billion in total annual AI consumption. Average monthly AI spend per firm grew 36% between 2024 and 2025 (CloudZero).

The inefficiency. BCG estimates 30% of AI spend is unmanaged inefficiency: oversized models, duplicated tools, workloads with no allocation to outcomes. That is a $4.5B addressable opportunity for a platform that can identify, quantify, and reduce this waste while providing the governance layer regulators require.

Our serviceable market. Within the $15B total, we target regulated mid-market enterprises in Europe and North America in insurance, financial services, utilities, healthcare, and manufacturing. These sectors represent ~30% of the addressable base, a SAM of $1.3-1.5B. They have the AI spend, the governance mandate, and the regulatory exposure that makes our solution not optional but necessary.

TAM

$15B

50k firms × $300K avg AI spend.

SAM

$1.5B

Regulated sectors, EU + North America

SOM

$10M ARR

100 clients × ~$100K ACV · Year 3. At this target ACV, each client represents 20-40% of directly addressable savings through cost optimisation alone, prior to compliance cost avoidance or productivity improvements

Why now — three converging deadlines

Aug 2, 2026 — **EU AI Act**

High-risk AI obligations live. Fines up to €15M or 3% of global turnover for non-compliance.

2025–2026 — **CSRD**

Corporate Sustainability Reporting Directive requires AI energy consumption disclosure for ~50,000 EU companies.

Ongoing — **Sector mandates**

EBA, EIOPA, HIPAA AI guidance tightening. Regulated industries face layered compliance demands.

VALUE PROPOSITION

What value do we create?

We tell you what AI is costing, creating, risking, and consuming — and what to change.

RightSize AI is the first vendor-neutral platform that gives leadership teams of regulated enterprises a consolidated, continuous view of their AI portfolio across four dimensions no existing tool addresses together.

The promise to the buyer: a CFO, CIO, or CRO opens RightSize AI on a Monday morning and sees what AI cost last week, what it delivered by use case, where there is risk, how the portfolio compares to sector peers, and three specific actions to improve the return curve. This is not a developer tool. It is a leadership instrument.

What makes it different: vendor-neutral — we sit above AWS, Azure, GCP, OpenAI, and Anthropic without preferring any. No new infrastructure required.

Vendor-neutral · No infrastructure change · Works with what you already have

Cost intelligence

Cost per inference, per use case, per team. Model quality gap: where an expensive frontier model is doing work a cheaper model could handle at equivalent quality. The FinOps Foundation calls this the 'LM Model Choice Quality Score.' Most companies have no idea this gap exists, let alone how large it is.

Value intelligence

Time saved per workflow, throughput improvements, error rate reductions — and the 'AI Tax': the share of time saved that is lost to rework and correction. Workday (2026) found this averages 37% across enterprises. Most organisations track AI adoption; almost none track whether adoption is net positive at the use-case level.

Risk intelligence

Shadow AI usage, model provenance (which vendors' models are touching which data), data exposure events, violations, and compliance status against internal policies, the EU AI Act and sector-specific frameworks. This gives the CRO and compliance team the audit trail they currently do not have.

Sustainability intelligence

AI energy consumption and carbon equivalent by workload and cloud region - using electricitymaps.com data to identify the least carbon-intensive regions. This produces the data required for CSRD reporting and EU AI Act transparency obligations. No current AI measurement tool offers this in a regulatory-ready format.

THE SOLUTION

How does it work?

Three layers — each delivering standalone value, each enabling the next

RightSize AI is structured as **three sequential product layers**, allowing non-AI native firms to bridge the structural expertise gap with modular, immediate value. Each tier deepens operational control and strategic acceleration, transforming unmanaged AI scaling from a systemic liability into a verifiable strategic asset.

Deployment is tiered to match organizational maturity: T**ier 1 provides foundational visibility** through zero-touch instrumentation, while **Tiers 2 and 3 deploy as an active API gateway to enable real-time prevention and workload orchestration**. By automating the complexity of routing, privacy, and cost-efficiency, we turn governance into an accelerator that makes the most responsible execution path also the fastest and most convenient choice for users and developers.

The **strategic moat** is built on proprietary, sector-wide benchmarks derived from multi-client execution data. This creates a powerful network effect: as more regulated firms join the platform, our data provides every client with unique fiduciary and competitive signals that cannot be replicated by internal teams or late-market entrants without access to the same sector-wide scale.

01 Enterprise Observability

MVP v1.0

30-day deployment · no infrastructure change

Connects to cloud billing APIs (AWS, Azure, GCP), model API logs (OpenAI, Anthropic, and others), SaaS usage data, and existing observability systems. Normalizes costs across all pricing models - per-token, per-API-call, per-seat, and reserved capacity. Produces a unified AI portfolio view: every tool, every use case, and every team, mapped to cost and usage. This is the '**Total Transparency**' layer: providing the fiduciary truth required to identify the structural resource leak and bridge the gap between technical execution and executive oversight.

02 Active Governance

v2.0

Cost vs. value · model quality gap · peer comparison

Deploys as a real-time API gateway and control layer sitting between developers and global model providers. Dynamically enforces organizational policies for security, privacy, and cost at the point of choice. Automatically identifies and blocks 'cost-creepers' and non-compliant geographic routing before execution occurs. This is the '**Immediate Prevention'** layer—providing a live shield that stops 'rogue scaling' and compliance violations before they become multi-billion dollar liabilities, moving the enterprise from 'after-the-fact' reporting to proactive risk management.

03 Intelligent Acceleration

v3.0

Recommendations · policy enforcement · compliance exports

Provides a unified 'One-Stop' marketplace for developers to subscribe to vetted AI services, removing the friction of manual vendor, security, and compliance reviews. Features a 'What-If' Simulator for predictive policy modeling, allowing leadership to simulate the financial and compliance impact of new AI usage before deployment. Automates complex workload orchestration by selecting the fastest, cheapest, and most compliant execution path for any AI usage—from real-time inference to large-scale model training. This is the 'Strategic Enablement' layer—transforming governance into an accelerator that makes the most responsible AI choice also the fastest and most convenient path for the engineering team.

COMPETITIVE LANDSCAPE

How is the solution differentiated?

The market has cost tools and productivity tools. It does not have a value intelligence platform.

CloudZero / Apptio:

Answer 'what did we spend?' with precision — but do not address value, risk, or benchmarking. Infrastructure cost tools, not AI governance platforms.

Larridin:
Measures AI utilisation and workforce productivity. Focus is on individual employee AI adoption, not the organisation's portfolio at use-case and outcome level. No model governance, compliance, or sustainability.

LiteLLM / PortKey / Martian:
Address model routing and cost optimisation for engineering teams. Developer-facing, not leadership dashboards. Answer a narrow version of the cost question; do not surface value, risk, or compliance.

Native cloud dashboards:

Provide billing data for their own services only. No consolidated view across multi-provider, multi-cloud environments, or visibility on shadow AI.

Big consulting (Deloitte, KPMG AI):
Can answer all four questions — through manual engagement, at 10–20x the cost, with no ongoing monitoring. A point-in-time deliverable, not a continuous intelligence layer.

	What does AI cost?	What does AI create?	Where is AI risky?	How do we compare?
CloudZero / Apptio	●	○	○	○
Larridin	○	●	○	○
LiteLLM / PortKey	●	○	○	○
Native cloud dashboards	●	○	○	○
Big consulting (manual)	●	●	●	○
RightSize AI	●	●	●	●

Our benchmarking database — cross-client, sector-segmented — accumulates value with each client. Once we have 20 insurers on the platform, each can see how their AI portfolio compares to direct peers. No internal team can build this; no new entrant can replicate it without years of data collection.

BUSINESS MODEL

What is the revenue model?

Annual subscription SaaS — priced for how regulated enterprises actually buy

Primary revenue: annual subscription. RightSize AI is sold as an annual contract, priced by the number of active AI use cases under management. Target ACV is $80,000–$100,000 for a mid-market client with 5–10 active AI deployments — deliberately positioned in the same bracket as enterprise GRC and compliance software, which our buyers already procure on annual contracts.

Secondary revenue: platform-led consulting. When the platform identifies optimisation opportunities — a use case with a 4x model quality gap, a shadow AI exposure, a compliance deadline — clients can engage RightSize for a structured implementation project. Triggered by the platform's own recommendations, scoped tightly. Expected to represent 20–30% of revenue by year two, at higher margins than the subscription tier.

Why subscription over usage-based. Our buyers — CFOs, CIOs, CROs — budget annually, need cost certainty, and procure compliance tools on fixed-fee contracts. A usage-based model introduces the same budget unpredictability we are selling to eliminate. Gross margin target: 80–85% on subscription, 40–50% on consulting; blended ~75% at scale.

Starter — **$50K /yr**
Up to 5 AI use cases

Growth — **$90K /yr**
Up to 15 AI use cases

Enterprise — **$150K+ /yr**
Unlimited + compliance exports

++ Platform-led consulting · 20–30% of revenue at scale · 40–50% gross margin

Revenue ramp ($M ARR)

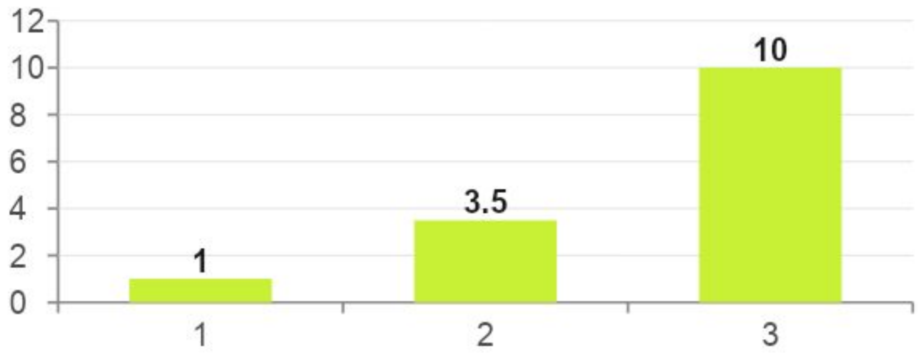

RISK FRAMEWORK

What are the key risks to be investigated and mitigated?

Risk Map: Pressure Points & Counter-Moves

HIGHEST ASSUMED RISK AT PRESENT	WHY IT MATTERS	MITIGATION
Enterprise sales cycles slower than forecast	Delays revenue ramp; increases burn	Founder-led sales into existing networks; EU AI Act deadlines compress procurement; low-friction 30-day visibility assessment as entry product
Integration complexity across fragmented AI stack	Slows deployment; risks failed pilots	Tiered deployment starting with observability only (no infra change); standardised connectors for AWS, Azure, OpenAI; expand coverage incrementally
Cloud providers replicate functionality	Hyperscalers could bundle similar tools at the infra layer	Vendor-neutral fiduciary layer; cross-client benchmarking moat cannot be replicated internally; CFO/CRO buyer, not developer tooling
Regulatory fragmentation and evolution	Requires continuous product adaptation across regions	Modular compliance architecture; focus on EU + regulated sectors; align outputs to reporting standards, not specific rules
Difficulty proving ROI to CFO	CFO won’t renew without measurable value; slows expansion	Embedded ROI tracking — cost savings, model optimisation, “AI Tax” measurement; benchmark comparisons; pilot-led value proof
Enterprise data trust and security concerns	Hesitation to expose sensitive AI usage data	Metadata-level analysis — no storage of sensitive payloads; enterprise-grade security and compliance certifications; on-prem option on roadmap

GO-TO-MARKET · FINANCIALS

What is our plan to reach the market?

Land with visibility. Expand with value. Scale through partners.

Phase 1 — Land (Year 1). Entry offer: 30-day AI portfolio visibility assessment — time-boxed, no infrastructure change. Target sectors: insurance and financial services first. The first 10 clients come from existing networks: Etay's enterprise investment relationships, Alex's industry contacts, Ravi's regulated financial institution network, and the TRIUM alumni network. Sales motion: CIO opens the door, CFO approves budget, CRO accelerates the timeline.

Phase 2 — Expand (Year 2). Once a client has the dashboard, the expansion path is natural: add use cases, add the risk and compliance module, engage the consulting tier for guided optimisation. The August 2026 EU AI Act deadline compresses procurement timelines by 1–2 months in European markets. Target net revenue retention: 115–120%.

Phase 3 — Scale (Year 3+). System integrators (Accenture, Capgemini, NTT Data) are delivering AI transformation at our target clients and need a governance layer to wrap around their implementations. White-label or referral partnerships targeting 30% of new logo revenue by year three.

Phase 1 · Land

Year 1

Direct outbound to **CIOs** and **CFOs** in **EU and US markets**. Entry: 30-day AI portfolio visibility assessment. First 10 clients from existing networks. Target: insurance and financial services.

10–12 clients · $1M ARR, Burn: 1.5-2m, Focus: Product Valuation

Phase 2 · Expand

Year 2

Deepen within existing clients — add use cases, compliance modules, consulting engagements. EU AI Act deadline compresses sales cycles. NRR target 115–120%.

40 clients · $3.5M ARR, Improved CAC via References, Channel Leverage

Phase 3 · Scale

Year 3

SI channel partnerships with Accenture, Capgemini, NTT Data. White-label and referral arrangements. Channel to contribute 30% of new logo revenue.

100 clients · $10M ARR, Path to EBITDA-positive scaling

THE TEAM

Why are we the right team?

We didn't find this problem in a research paper. We lived it.

Every member of this team has encountered the problem RightSize AI is solving — not read about it, but encountered it professionally in systems where the cost of getting AI wrong was material. This is not a team built for a project. This is a team built for this problem.

The closing argument.
88% of organisations use AI. 95% of pilots fail to deliver measurable return. The EU AI Act compliance clock is now running. The question is not whether enterprises need RightSize AI. The question is how quickly we can get in front of them — and we have the relationships, the domain knowledge, and the technical credibility to do so.

Key milestones from now.

May 2026: Prospectus submitted. 5-10 customer discovery interviews being conducted.

July 2026 (Module 4): Business plan stress-tested. First pilot agreement targeted.

Nov 2026 (Module 5): Written report in draft. MVP technical specification complete.

GE

Guillaume Ehinger

Google

Product & Technology

Ran AI systems where performance, cost, and security all matter simultaneously at planetary scale. Has seen the model-selection and cost-attribution problem from inside one of the world's largest AI deployments.

EZ

Etay Zwick

Attuned Futures

Market Strategy

Connects frontier technology to real enterprise investment decisions. Has the buyer relationships in the CFO and CIO community that will produce our first clients, and the pattern recognition to know when a product is genuinely needed.

AP

Alex Paterson

Cisco Systems

Monetization & Sales

Turns complex AI and network systems into measurable, controllable operations for large enterprises. Understands the enterprise sales motion, multi-stakeholder buying, and how to build a pipeline in regulated industries.

DM

David Mindham

EDP Renewables

Operations & Ecosystem

Priced risk and operational constraints at a heavily regulated energy company where inefficiency is directly costly and compliance is non-negotiable. Brings the governance and regulatory lens that makes our product credible to the CRO audience.

RG

Ravi Govender

Momentum Group

Financials & Funding

Built governed, scalable AI systems inside complex financial institutions. Focus on shaping pricing strategy, capital pathway, enterprise fiduciary logic and investor readiness while bringing direct experience as both product builder and archetypical enterprise buyer

The market is ready. The team is ready. Let's build it.

APPENDIX

What follows — supporting frameworks

Appendix: Strategic Frameworks

A1

BUYER FIT

Value Proposition Canvas

Match between RightSize's value map and the CFO/CIO/CRO buyer's jobs, pains and gains.

A2

OPERATING MODEL

Business Model Canvas

Nine-block view of partners, activities, resources, channels and unit economics.

A3

MODEL COHERENCE

Odyssey 3.14 Assessment

Value Proposition → Architecture → Profit Equation: does the business work as a system?

A4

TAILWIND

Regulation as Tailwind

Why EU AI Act timing converts compliance pressure into commercial demand for RightSize.

These appendices test the core argument from four angles — buyer, model, system, and timing.

VALUE PROPOSITION

Hypothesised — pre-validation

Value Proposition Canvas

TARGET **CFOs, CIOs & CROs at regulated mid-market enterprises (FS, insurance, healthcare)** · Economic buyer: CFO · Technical gatekeeper: CIO/CTO · Risk accelerator: CRO

VALUE PROPOSITION

GAIN CREATORS

- Peer benchmarking
- Predictive optimisation
- EU AI Act-aligned outputs
- Board-ready reporting

PAIN RELIEVERS

- Unified cost tracking
- Audit trails & documentation
- Anomaly detection on usage
- Compliance evidence

PRODUCTS & SERVICES

- Observability platform
- Governance & compliance layer
- Benchmarking engine
- AskAI natural-language interface
- Risk & ROI dashboards

FIT

CUSTOMER PROFILE

GAINS

Visibility + ▪
compliance proof
Quantified ROI
Strategic edge

CUSTOMER JOBS

- Allocate AI cost
- Ensure compliance
- Stay in control
- Be seen as disciplined
- Brief the board credibly

PAINS

No cost visibility
Board ROI pressure
Fines €15M / 3%
Model leakage

RightSize AI enables CFOs and risk leaders to understand, optimise, and govern enterprise AI usage in real time.

BUSINESS MODEL

Hypothesised — pre-validation

Business Model Canvas

KEY PARTNERS

- Cloud providers (AWS/Azure/GCP)
- Foundation model providers
- Systems integrators
- Compliance & audit firms

KEY ACTIVITIES

- Product build & data aggregation
- Sales · compliance mapping

KEY RESOURCES

- Benchmarking dataset (critical)
- Engineering · regulatory IP
- Enterprise relationships

VALUE PROPOSITIONS

CFO
- Cost + ROI visibility

CIO
- Control + architecture

CRO
- Compliance + audit

CUSTOMER RELATIONSHIPS

- High-touch onboarding
- Land → expand → scale

CHANNELS

- Direct (founders, outbound)
- Indirect (SIs, partners)
- Product-led (platform)

CUSTOMER SEGMENTS

PRIMARY
- Regulated mid-market
- FS · insurance · healthcare

SECONDARY
- Large enterprises via SIs

COST STRUCTURE

FIXED
- Engineering · sales · compliance

VARIABLE
- Cloud · delivery · partner fees

REVENUE STREAMS

PRIMARY
- SaaS subscription ·$50–150K ACV

SECONDARY
- Consulting (20–30% of revenue)

FUTURE
- Benchmarking data · premium modules

ECONOMIC MODEL · hypothesised CFO layer · the deal logic

ACV	GROSS MARGIN	CAC	PAYBACK	NRR
$90–100K	80–85%	$75–125K	12–18 mo	115–120%

VALUE CAPTURE LOGIC

Client spends **$500K** on AI → **$150K waste** → **$75K recoverable** → **$90K fee** justified

BUSINESS MODEL · ODYSSEY 3.14

WHAT → HOW → WHY IT WORKS

Odyssey 3.14 Business Model Assessment

01 · WHAT

Value Proposition

REDUCE COST
- Eliminate 20–30% AI waste
- Optimise model selection

REDUCE HASSLE
- Single view across AI usage
- Automated compliance

EXPAND BUYER
- CFO / CRO, not just AI teams

ADD INTELLIGENCE
- AskAI + benchmarking + recs

NETWORK EFFECT
- Cross-client benchmarking dataset

SUMMARY
Vendor-neutral AI portfolio intelligence — cost, compliance, decisions.

DRIVES

02 · HOW

Value Architecture

3-LAYER PLATFORM
- Observe → govern → optimise
- Vendor-neutral control layer

KEY ACTIVITIES
- Multi-cloud data ingestion
- Cost normalisation · policy enforcement

STRATEGIC MOAT
- Benchmarking dataset (compounds)
- Regulatory mapping engine

PARTNERSHIPS
- Cloud (AWS, Azure) · Models (OpenAI, Anthropic) · SIs

SUMMARY
Lightweight control + intelligence layer above fragmented AI infrastructure.

DRIVES

03 · WHY IT WORKS

Profit Equation

REVENUE
- SaaS $50–150K ACV · consulting 20–30%
- Future: benchmarking data, modules

COST
- Eng + sales fixed · cloud variable
- CAC heavy early, then leverages

CAPITAL EFFICIENCY
- Cloud-native · low working capital
- 80–85% gross margin

VALUE CAPTURE
- Client saves $75–100K → fee $90K
- Self-funded ROI

ROCE
- High margin × high capital turnover

SUMMARY
High-margin SaaS · low capital intensity · ROI-aligned, regulation-pulled.

INSIGHT

RightSize AI transforms AI governance from a cost centre into a high-margin, capital-efficient SaaS model — driven by regulatory demand and measurable economic value.

REGULATORY LANDSCAPE

What is the regulatory environment?

Regulation is not a headwind — it is our tailwind

Rather than framing regulation as a risk to manage, RightSize AI treats the regulatory environment as a primary demand driver. Three regulatory developments are creating mandatory requirements that our platform is purpose-built to fulfil.

Intellectual property and defensibility. Our primary IP is the benchmarking database — cross-client, anonymised, sector-segmented AI performance data that accumulates value with each new client. Not replicable by a new entrant without years of client acquisition. Secondary IP is our regulatory compliance reporting methodology: the mapping between RightSize outputs and the specific documentation requirements of the EU AI Act, CSRD, and sector frameworks.

Target sectors. Insurance, financial services, utilities, healthcare, and manufacturing — each with layered compliance demands beyond the general EU AI Act and CSRD, including Solvency II / EIOPA for insurance, EBA AI governance guidance for banks, and evolving HIPAA requirements for healthcare.

The EU Commission opened a consultation in April 2026 on developing a framework for measuring AI energy consumption under the AI Act — the regulatory machinery is actively being built around the exact problem we solve.

EU AI Act

August 2, 2026 The world's first comprehensive AI regulation classifies AI systems by risk level. For regulated enterprises in financial services, insurance, and healthcare — many AI systems will be high-risk under Annex III. Requirements include: technical documentation, human oversight mechanisms, decision logging for auditability, and energy consumption disclosure. Penalties: up to €15M or 3% of global annual turnover.

RightSize provides: audit trails, model documentation, compliance exports, energy reporting

CSRD

2025–2026 reporting cycle Corporate Sustainability Reporting Directive requires large companies and listed SMEs to report on sustainability impacts including AI-related energy consumption. Applies to approximately 50,000 companies in the EU. RightSize's sustainability module measures AI energy consumption per workload and cloud region, producing data compatible with European Sustainability Reporting Standards.

RightSize provides: energy per workload, carbon equivalent by region, CSRD-ready exports

Sector frameworks

Ongoing, tightening Our target verticals have their own AI governance obligations layered on top of general regulation. Solvency II and EIOPA guidelines for insurance require model risk management and documentation. EBA AI governance guidance for banks requires explainability and audit trails. HIPAA's evolving AI guidance for healthcare requires risk assessments and data handling controls.

RightSize provides: model provenance, risk flags, sector-specific reporting outputs